


Optimization and Gradient Descent

Optimization in the context of machine learning


- Given some sample output points y_i for $i = 1, \dots, N$ and some sample input points x_i for $i = 1, \dots, N$ and the function $f()$ of the form $f(x_i, w_k)$, for $k = 1, \dots, K$ where K is the number of parameters. Find w_k for $k = 1, \dots, K$ that make $f(x_i, w_k)$ as close as possible to y_i for all $i = 1, \dots, N$
- $j(w_k) = \sum_{i=1}^N |(y_i - f(x_i, w_k))|^2$
- *How to find w_k that minimize this cost function $j(w_k)$*

Finding minia of $j(w_k)$



The derivative is *zero* at any local maximum or minimum.

Finding minima of $j(w_k)$



Assumption

- Let assume the cost function is $f(x)$ where x is some parameters.
- How to find the best x that minimizes the cost function?

The derivative is *zero* at any local maximum or minimum.

One way to find a minimum: set $f'(x)=0$ and solve for x .

$$f(x) = x^2$$

$$f'(x) = 2x$$

$$f'(x) = 0 \text{ when } x = 0, \text{ so minimum at } x = 0$$


Rate of change in derivative

The slope at a point is called the **derivative** at that point

Intuition: Measure the slope between two points that are really close together

$$\frac{f(x + c) - f(x)}{c}$$

Limit as c goes to zero



The derivative is *zero* at any local maximum or minimum.

One way to find a minimum: set $f'(x)=0$ and solve for x .

- For most functions, there isn't a way to solve this.
- Instead: algorithmically search different values of x until you find one that results in a gradient near 0.


Finding minima

If the derivative is positive, the function is **increasing**.

- Don't move in that direction, because you'll be moving away from a trough.

If the derivative is negative, the function is **decreasing**.

- Keep going, since you're getting closer to a trough




$$f'(-1) = -2$$

At $x=-1$, the function is decreasing as x gets larger. This is what we want, so let's make x larger.

Increase x by the size of the gradient:

$$-1 + 2 = 1$$




$$f'(-1) = -2$$

At $x=-1$, the function is decreasing as x gets larger. This is what we want, so let's make x larger.

Increase x by the size of the gradient:

$$-1 + 2 = 1$$




$$f'(1) = 2$$

At $x=1$, the function is increasing as x gets larger. This is not what we want, so let's make x smaller.

Decrease x by the size of the gradient:

$$1 - 2 = -1$$

Learning rate solves this jumping behavior



$$f'(-1) = -2$$

$$x = -1 + 2(.25) = -0.5$$

$$f'(-0.5) = -1$$

$$x = -0.5 + 1(.25) = -0.25$$

$$f'(-0.25) = -0.5$$

$$x = -0.25 + 0.5(.25) = -0.125$$

Eventually we'll reach $x=0$.

Learning Rate

In order to guarantee that the algorithm will converge, the learning rate should decrease over time. Here is a general formula.

At iteration t :

$$\eta_t = c_1 / (t^a + c_2),$$

where $0.5 < a < 2$

$$c_1 > 0$$

$$c_2 \geq 0$$

Gradient Descent

Gradient descent is guaranteed to eventually find a *local* minimum if:

- the learning rate is decreased appropriately;
- a finite local minimum exists (i.e., the function doesn't keep decreasing forever).

Gradient Descent

1. Initialize the parameters \mathbf{w} to some guess (usually all zeros, or random values)
2. Update the parameters:
$$\mathbf{w} = \mathbf{w} - \eta \nabla L(\mathbf{w})$$
$$\eta = c_1 / (t^a + c_2)$$
3. Repeat step 2 until $\|\nabla L(\mathbf{w})\| < \theta$ or until the maximum number of iterations is reached.

Stopping Criteria

For most functions, you probably won't get the gradient to be exactly equal to **0** in a reasonable amount of time.

Once the gradient is sufficiently close to **0**, stop trying to minimize further.

How do we measure how close a gradient is to **0**?

Stopping Criteria

Stop when the norm of the gradient is below some threshold, θ :

$$\|\nabla L(\mathbf{w})\| < \theta$$

Common values of θ are around .01, but if it is taking too long, you can make the threshold larger.

Description of Gradient Descent Method

- Algorithm (Gradient Descent Method)
 - **given** a starting point $x \in \text{dom } f$
 - **repeat**
 1. $\Delta x := -\nabla f(x)$
 2. Line search: Choose step size η via exact or backtracking line search
 3. Update $x := x + \eta \Delta x$
 - **until** stopping criterion is satisfied
- Stopping criterion usually $\|\nabla f(x)\|_2 \leq \epsilon$
- Very simple, but often very slow; rarely used in practice